

Build the organisation.

A deliberately engineered recurrent mind: attention, imagination, memory, intention and self-observation.

36 REAL COMPLETIONS


Two engineering runs with DeepSeek. 31 structured results applied; five rejected results retained.

8 CONTINUING EPISODES

The corrected run completed its full trajectory. One chosen intention survived the interruptions and handler reconstruction.

2 IDEAS RETURNED

Images generated in earlier episodes later became the selected focus. Their recorded IDs establish the connection.




What this establishes

An executable recurrent prototype whose internal records affect later operation. Accurate self-knowledge remains incomplete. Subjective experience has not been established.

The research objective is construction of a mind. Office-task efficiency is not used as a substitute for that objective.

One loop, shared consequences.



Implementation map. Arrows show information flow through software components, not neural activity.



The operator can interrupt, pause, inspect or resume. The scheduler continues discrete ticks after the browser closes when enabled. Code-revision proposals do not install themselves.

Generated content came back into play.

R2: eight successive episodes from the same persistent agent, including three operator interruptions.



EPISODE 1 > EPISODE 3

"River cutting its own channel" was imagined, saved and later selected by its persistent object ID.


EPISODE 2 > EPISODE 5

"Two rooms connected by a corridor" returned after intervening episodes. The controller linked imagination to attention.

Remove a connection. Observe the result.

Five conditions, two continuations each, from one shared R2 checkpoint. These are component-route interventions.

Condition	Observed focus	Runs
Full	Imagination	2/2
Affect removed	Tension	2/2
Focus replaced	Inquiry	2/2
Memory removed	Imagination	2/2
Self-model removed	Imagination	2/2



CAUSAL CONNECTION DEMONSTRATED

Affect removal changed the selected focus in both repetitions.
Workspace replacement changes it by construction.

IDENTITY EFFECT NOT ESTABLISHED


The self-model projection removal did not change focus. Other self-relevant context remained. No identity-specific inference follows.

Its story about itself was wrong.

Episode 8 supplied the most informative failure: confident self-history contradicted the recorded events.

Model claim: "ticks 3,5,7 were reflections"

Recorded events: 3 = imagination; 5 = imagination; 7 = tension.



Mean Brier: 0.886

Uniform baseline: 0.857. On four eligible autonomous ticks, the model did not beat this simple forecast. The sequence is too small and dependent for a population claim.

Engineering response

The final build supplies actual event history and validates tick/kind references. This correction has local test coverage; the additional live validation was blocked.

Talk to the same continuing system.

Actual R2 records shown in the new interface. Read-only evidence preview; interactive controls were tested separately with synthetic responses.

Think one cycle

Explore

6 cycles

Begin

Pause autonomous thought

Ready

CURRENT WORKSPACE

Compare the previous prediction of attention with what occurred; revise the working self-model if needed.

reflection · engine_monitor · 3 retrieved memories

reflection		0.98
reflection		0.63
reflection		0.63
reflection		0.59
reflection		0.59
imagination		0.57

Each focus feeds cognition, memory, drive satisfaction and the next

Talk with the mind

misattribution without manufacturing it. That may be the closest permissible test of whether recorded provenance constrains revision differently than mere context. I will refine the intention toward designing a protocol around naturally occurring provenance anomalies rather than artificial edits.

Cycle 8 · affective_integration

Reviewing prediction error: I predicted reflection was unlikely (0.1) but reflection occurred, yielding Brier 0.955. The uniform prior I used was uninformative; a better predictor should weight recent history (ticks 3,5,7 were reflections) and the engine's self-review tendency. I should revise my prediction policy toward continuity rather than uniform spreads. No belief changes warranted; this is a calibration update.

Message

The displayed episode-eight error is retained verbatim. An inspectable mistake is part of the experimental record.

Give self-observation a factual substrate.

The next construction target is a learned causal model of its own organisation and possible interventions.

ALREADY BUILT

- Shared attention
- Persisted intentions
- Imagination re-entry
- Appraisal-to-attention coupling
- Memory retrieval
- Forecast instrumentation
- Conversation and pause

ADDED AFTER THE FAILURE

Eight recorded attention events in context. Structured event references checked before state update. Existing intention priorities can now alter future bids.

NEXT EXPERIMENT

Predict what removing a component will do. Compare own event history, matched third-person history and withheld history across independently started agents.

What remains unsettled

A working recurrent organisation does not settle whether there is subjective experience. It also does not establish calibrated self-knowledge: this run contains a clear counterexample. The programme can now investigate these questions through changes to real components.

No original A-F gate is newly awarded by this component diagnostic. No weight update or autonomous code promotion occurred.

A complete package, with the failures included.

The ZIP contains the Functions, UI, tests, report, analysis and collected V8 evidence. Production deployment was not performed in this pass.

RUN AND INSPECT

```
npm ci
npm test
node scripts/live-mind.mjs new-run
node scripts/analyse-mind.mjs analysis new-run
```

Set DEEPSEEK_API_KEY in the environment. The full schedule uses 18 calls. Existing output directories are refused.

DEPLOY THE COMPLETE APP

Unzip; select the existing site with Netlify CLI. Set DEEPSEEK_API_KEY and SOVEREIGNTY_ACCESS_TOKEN for Functions.

```
npx netlify-cli deploy --build --prod
```

Source-folder drag and drop alone does not install Functions.

EVIDENCE ACCOUNTING

36 confirmed completions; 311,279 tokens. R1: 13 applied, five rejected. R2: 18 applied, two with explicit brace repairs. 52 tests pass. The six-call follow-up was blocked by automatic approval review before any completion; no retry was made.

AUTONOMY AND PROVENANCE

```
AUTONOMY_ENABLED=true
AUTONOMY_INTERVAL_MINUTES=15
AUTONOMY_DAILY_CALL_LIMIT=24
```

Existing environment values override defaults. Raw transcripts, checksums, commands and limitations accompany the report.

Netlify reference: docs.netlify.com/build/functions/get-started/ | [scheduled-functions/](#) (accessed 09 Sep 2026)