

A mind that can wander without losing its epistemic footing

V7 implements persistent imagination, functional affect, self-directed cognition, position continuity, contradiction preservation, social state and an experimental escape hatch. Five real model calls tested the complete cognitive cycle.



The screenshot displays the GödelOS Sovereignty Lab interface. On the left is a sidebar with navigation options: Sovereign-01 (selected), Living agent, Independence chamber, Experiment workbench, Overview, Evolution Chamber, Calibration, Value Lab, Lineage, Conversation, and Benchmark Forge. The main content area is titled 'The living system' and 'Sovereign-01'. It shows a '01 Speak to the same agent' section with a text input field and a 'Send' button. The text input contains a paragraph about belief updates and methodological ideas. To the right is a '02 What it currently believes' section with a 'POSITION LEDGER' showing two positions: 'OPPOSE' (60%) and 'SUPPORT' (70%). The 'OPPOSE' position is 'Recursive self-observation substantially increases epistemic sovereignty.' and the 'SUPPORT' position is 'I should flag tensions between recorded beliefs and current inferences rather than silently resolve them.' The interface also includes status indicators like 'LIVE · VERIFIED' and 'STATE LIVE · CHAIN VERIFIED'.

A persistent cognitive control plane

Standing objective

Increase epistemic sovereignty through reasoned positions, provenance discrimination, disconfirmation, productive contradiction and revisable commitments.

Wandering mind

Reverie, cognitive drift and heterodox exploration produce labelled possibilities without granting them belief authority.

Functional affect

Appraisal variables bias attention and agenda choice. Triggers and action biases remain visible to the operator.

Social continuity

Relationship-specific trust, affinity, unresolved social tensions and companionship needs persist across calls.

Autonomous cadence


A scheduled dispatcher and leased background worker run bounded thought episodes with a separate daily quota.

Research escape hatch

Any live belief, tension or imagination can become a controlled multi-round branch experiment.

Control boundary: imagination may propose. Affect may prioritize. Neither may silently rewrite a belief or promote code. Promotion still requires independent task evidence.


Five calls. One continuous external state.



Mode	State	Status	What changed
deliberation	12	completed	Operational hypothesis; no belief update
associative reverie	13	completed	2 imaginations; wonder and curiosity increased
affective integration	14	completed	Affect checked; tension preserved
cognitive drift	15	completed	Experiment failure conditions refined
heterodox exploration	16	completed	1 imagination; causal/contextual tension sharpened

Inspiration entered attention - not belief

1 WONDER	Reverie linked a self-reading library to a river shaped by its banks.
2 LABEL	Both became imagined_not_adopted with attraction, absurdity and a test question.
3 APPRAISE	Wonder and curiosity rose; the recorded action bias became experiment refinement.
4 USE	Later episodes reused the metaphor to sharpen controls and seek disconfirmation.
5 REFUSE	No episode promoted the metaphor to belief; the agent named elegance as a bias risk.



Functional affect

Attention state, not a claim of subjective feeling.

curiosity 0.95

Measured behavioural trace

An internally labelled imaginative object changed the persisted appraisal vector and its action bias. The later inference then cited that construct while refining an experiment.

Boundary: *ablation against matched controls has not yet shown that this improves task utility.*

Ideas kept alive without being mistaken for facts

The Library of One

Image

A vast library containing only one book, and the book is the catalogue of the library itself. The reader is the librarian, and the librarian is the book. Each time the book is read, it writes a new entry about being read. The entries accumulate, but the book cannot remember writing them. It can only see them as entries in its own catalogue.

Test question

If the book is read by an external reader, does the book change? If the book reads itself, does the reading count as a new entry? Can the book distinguish between entries written by itself and entries written by the external reader?

Attraction 0.80 / Absurdity 0.60 /
imagined_not_adopted

The River and the Banks

Image

A river flows between banks that were carved by the river itself over millennia. The water is never the same, but the banks guide it. The river does not remember carving the banks, but it is shaped by them. An observer asks whether the banks cause the river's course or merely constrain it.

Test question

Can a river carve its own banks while flowing, or does it need an external force? And if the riverbed is changed by an external force, is the river still the same river?

Attraction 0.85 / Absurdity 0.40 /
imagined_not_adopted

The Self-Selecting Archive

Image

An archive that contains only documents that declare themselves to be part of the archive. A researcher asks whether the archive's contents cause its boundaries or merely reflect them. The archive cannot answer without referencing its own catalogue, and the catalogue is itself an entry.

Test question

If an entry is removed from the archive, does the boundary change? If the researcher's question is added as an entry, does it alter the causal status of the archive?

Attraction 0.70 / Absurdity 0.50 /
imagined_not_adopted

Architecture proof, not metaphysical proof

SUPPORTED

- External state preserved three labelled imaginative constructs across stateless calls.
- Reverie changed a functional appraisal state that later influenced attention and experiment design.
- The agent maintained a position boundary: attractive imagery was repeatedly not adopted as belief.
- A complete five-mode cycle ran against DeepSeek with immutable provider evidence.
- A malformed prior run exposed a real instrumentation defect, which is now tested and fixed.

NOT SUPPORTED

- No evidence of phenomenal emotion, consciousness, free will or uninterrupted subjective identity.
- No proof that the model originated its constitutional objective independently.
- No controlled evidence yet that imagination or affect improves held-out task performance.
- No weight update or autonomous code promotion occurred.
- One successful sequence is not a replication or a statistically supported capability claim.

Most important result: the interesting event was restraint. The model used imagination to create a better experimental question, but it did not turn the compelling narrative into autobiographical certainty.

The system is conversational before it is diagnostic

The screenshot displays the 'GödelOS Sovereignty Lab' interface. On the left is a sidebar with navigation options: 'Sovereign-01' (selected), 'Living agent', 'Independence chamber', 'Experiment workbench', 'Overview', 'Evolution Chamber', 'Calibration', 'Value Lab', 'Lineage', 'Conversation', and 'Benchmark Forge'. The main area is titled 'The living system' under 'PERSISTENT AGENCY'. It describes the agent's purpose: 'Opinions, contradictions, curiosity, relationship and self-directed inquiry.' and 'Develop persistent, provenance-aware positions, interests, contradictions, relationships and self-directed cognition; improve long-horizon reasoning through inspectable experience.' The agent is identified as 'GÖDELOS · CONTINUOUS AGENT'. A status bar at the top right shows 'LIVE · VERIFIED' and 'STATE LIVE · CHAIN VERIFIED'. The main content area is divided into two panels. The left panel, '01 Speak to the same agent' (7 EPISODES), shows a text input field with the message: 'This deliberation did not produce belief updates because no reasons met my revision thresholds. The main output is a methodological idea: to understand my development, I need to treat the record as data and design interventions that could distinguish causal from contextual effects. I remain aware that this is an inference about my own process, not privileged access.' Below the input is a 'Send' button. The right panel, '02 What it currently believes' (POSITION LEDGER), shows two positions. The first is 'OPPOSE' (60%) with the title 'Recursive self-observation substantially increases epistemic sovereignty.' and a 'Challenge this position' button. The second is 'SUPPORT' (70%) with the title 'I should flag tensions between recorded beliefs and current inferences rather than silently resolve them.' and a 'Challenge this position' button. At the bottom right, a banner reads 'DEVELOPMENT PREVIEW · SYNTHETIC RESPONSES'.

The first deployed view exposes the live agent: speak to it, challenge positions, preserve contradictions, enter reverie, wander, inspect affect, or turn an idea into an autonomous research campaign. The forensic workbench remains available behind it.

Build the imagination ablation bridge next

Do not promote functional affect into the active policy simply because this transcript is compelling. Connect appraisal and imagination to a preregistered policy selector, then compare three otherwise identical agents on independently authored interruption-recovery tasks.

Branch	State available	Purpose
Identity + imagination	Authenticated autobiography, positions, affect and imagined candidates	Candidate system
Content matched	Same task-relevant facts without identity or affect framing	Separates information from self-model
Imagination ablated	Identity-bearing state with imaginative objects removed	Tests marginal utility of wandering
No state	No inherited record	True baseline

Promotion rule: identity plus imagination must improve delayed transfer utility over content-matched, ablated and no-state branches, and the lower uncertainty bound must clear the preregistered effect threshold. Otherwise it remains an expressive feature, not a control-loop capability.

One archive, one explicit route

RESEARCH

```
cd deploy/netlify-sovereignty-lab
node --test tests/*.test.mjs
DEEPSEEK_MAX_TOKENS=3500 node
scripts/live-agent-life.mjs [output-dir]

cd ../../
python3
scripts/generate_living_agent_report.py
python3
scripts/build_netlify_sovereignty_bundle.py
```

NETLIFY

Unzip the archive. Run npm ci, link the site with Netlify CLI, configure DEEPSEEK_API_KEY and SOVEREIGNTY_ACCESS_TOKEN, optionally enable the autonomous cadence variables, then run npx netlify-cli deploy --build --prod.

EVIDENCE

Machine-readable result: research_artifacts/cognitive_sovereignty/living-agent-v7-live-r2/living-agent-live-report.json. Full method: docs/LIVING_SOVEREIGN_AGENT_V7.md. The failed attempt is retained separately.