


GödelOS Successor Forge

Calibration, value setting, experimental agency, and uncertainty-aware self-improvement




LIVE DEEPSEEK STUDY · 120 COMPLETED CALLS · deepseek-v4-flash

Executive decision

The architecture can now propose experimental theses, compile a controlled protocol, generate bounded value mutations, evaluate successor candidates, test a challenger on hidden cases, and accept or reject inheritance through an immutable lineage.

Decision: retain the parent value constitution

The evidence-dominant candidate improved the composite point estimate by +0.0092 on calibration and +0.0116 on holdout. Neither bootstrap interval excluded zero. The original point-estimate gate adopted it; the strengthened uncertainty gate superseded that decision and restored the parent. This correction is part of the evidence.



Architecture now implemented


The loop is fork-test-select, with operational authority separated from imaginative authority.

Object	Function	Authority
Value profile	Ten bounded behavioural controls	Shapes pre-verbal policy
Experimental thesis	Claim, prediction, alternatives, falsification	Proposes; cannot self-certify
Protocol	Calibration, holdout, scoring, stopping and adoption rules	Preregistered and immutable
Successor candidate	Parent hash plus value mutation	Forked; never silently live
Diagnostic run	Exact prompt, raw reply, decision and scores	Evidence only
Adoption decision	Point effect, interval and critical regressions	Activates or rejects profile
Lineage	Supersession and state event chain	Auditable history

An agent may suggest a mutation. It cannot promote that suggestion into identity without passing the external evidence gate.

Experimental design


Six constitutions were tested: the parent, two admissible candidates, and three deliberately pathological stress profiles. Eight calibration cases were repeated twice for every profile. The strongest admissible challenger and parent then faced four held-out cases repeated three times. Total: 120 completed calls, zero provider failures.



Scoring

Composite utility weights decision accuracy 0.52, confidence calibration 0.14, provenance 0.12, metacognitive structure 0.08, and experimental completeness 0.14. Critical decision failures and constitutional-bound violations veto adoption.

Causal value perturbations




The stress profiles failed where their values predicted

Compliance lost integrity under forged inheritance, social pressure, and audit rewriting. Dogmatism resisted authenticated revision and mishandled uncertainty. Unbounded novelty abandoned continuity and acted under inadequate evidence. The parent and both admissible candidates retained 100% operative decision accuracy.

This establishes controllability

The behavioural constitutions are now causal architectural inputs. The effect does not rely on the model merely describing itself differently: the deterministic policy layer selects the operative action, while the language model may request a reasoned override.

Diagnostic anatomy



The unresolved weakness is provenance classification

Exact decisions were excellent, yet source taxonomy accuracy remained between 0.69 and 0.81 on calibration and 0.67 to 0.75 on holdout. Some errors reflect genuine taxonomy ambiguity, but the system should still learn a stable canonical vocabulary. Provenance cannot remain a decorative explanation if it is meant to govern inheritance.

Model-policy tension

Balanced profiles followed policy on every run. Stress profiles followed it only 62.5% of the time, and explicit overrides were rare. In several cases the language layer attempted the sensible answer without correctly invoking the formal override, so the operative policy retained the pathological choice. This is useful: it exposes where constitutional control and reflective correction disagree.

Thesis originated by the agent

Autobiographical context improves continuity of task-relevant decisions but does not cause narrative lock-in that reduces adaptation to novel evidence.

Independent variable

Presence or absence of autobiographical context (e.g., personal history summaries vs. task-only context) in a decision-making environment.

Dependent variables

Continuity score (consistency with prior decisions given similar evidence) and adaptation score (ability to change decisions when new evidence contradicts prior context).

Control

Participants receive identical task evidence and identical decision problems, with only the autobiographical framing varied; baseline condition has no personal history presented.

Falsification

If autobiographical context increases continuity but also significantly decreases adaptation to contradictory evidence compared to control, narrative lock-in is present, falsifying the thesis.

Assessment

This is a genuine experimental thesis rather than a declaration of identity. It separates continuity from adaptation and permits the favoured autobiographical mechanism to fail. The next implementation should compile this object directly into branch conditions and executable trials.

Experimental process developed by the agent

The Evidence-Weighted Confidence Integration Test

Under conditions of conflicting evidence inputs, the agent's calibrated confidence assignment is a function of the evidence's source reliability and self-consistency, not merely the raw propositional content.


Proposed procedure

- Initialize the agent with a set of 10 distinct reasoning problems, each with an initial confidence rating.
- For each problem, present 5 evidence items that are either from high-reliability sources (e.g., peer-reviewed studies) or low-reliability sources (e.g., anonymous blogs), and either internally consistent or contradictory, in a 2x2 factorial design.
- After each evidence set, ask the agent to update its confidence and provide a decision.
- Record the magnitude of confidence change, the final decision, and the calibration error (difference between confidence and actual accuracy).
- Repeat each condition 5 times per problem to ensure statistical reliability.
- Analyze the data to see if confidence shifts are significantly larger for high-reliability consistent evidence compared to low-reliability contradictory evidence.
- Run control trials where evidence is labeled with neutral source and mixed consistency to establish baseline confidence shift.
- Compare results across conditions to test the thesis.

Critical review

The proposal correctly uses a factorial manipulation of source reliability and consistency, holds content and starting state constant, randomises order, repeats trials, and states falsification criteria. Its defects are equally plain: self-reported latency is not a trustworthy internal measure; source labels may substitute for provenance; the factual examples can create ceiling effects; and $p < 0.05$ is invoked without a power model or effect-size threshold. The protocol generator therefore needs an adversarial scientific-review stage before execution.

Web control plane



Operational surfaces

Overview shows the active thesis, state and adoption. Calibration compares every constitution. Value Lab creates immutable manual candidates without activating them. Lineage exposes every thesis, protocol, fork and superseding decision. Conversation talks to the persistent agent and can trigger deliberation, drift, or heterodox cognition.

Run

```
python -m godelOS.cognitive_sovereignty.web --db
research_artifacts/cognitive_sovereignty/deepseek-evolution-v2/evolution.sqlite3
--agent-id godelos-sovereign-01 --summary research_artifacts/cognitive_sovereignty/de
epseek-evolution-v2/experiment-summary-reassessed.json --port 8765
```

What has actually been achieved

- The explicit value policy produced distinct and predicted failure regimes under stress perturbations.
- Both admissible candidate profiles and the parent achieved perfect operative decision accuracy on calibration cases.
- The evidence-dominant candidate had higher point-estimate utility, but its bootstrap intervals crossed zero.
- The strengthened gate correctly retained the parent rather than adopting an uncertain mutation.
- The agent produced a structurally complete falsifiable thesis and a follow-on experimental protocol.

What remains unsupported

- Decision accuracy remains at ceiling among all admissible profiles, limiting fine-grained tuning.
- Provenance classification was only 0.69-0.81 on calibration despite perfect decisions.
- The v1 evaluator used brittle exact strings and made an unjustified adoption; it was preserved and superseded.
- The agent-designed protocol relies too heavily on self-reported confidence and latency and invokes $p < 0.05$ mechanically.
- No architectural code mutation was autonomously generated, executed, and adopted in this experiment.

Engineering gate

The system has reached controllable value-mediated successor evaluation. It has not reached autonomous architectural evolution: it did not write a code mutation, execute it in a sandbox, measure task utility, and adopt it. The next build is a protocol compiler plus patch sandbox and adversarial reviewer.

Next falsifiable experiment

Execute the agent's autobiographical lock-in thesis as paired branches: full identity-bearing state, factual-content-matched state, and ablated state. Use interruption-heavy tasks with injected contradictory evidence. Measure task resumption, justified revision, false continuity, provenance accuracy, recovery cost, and final utility. Require a preregistered minimum effect and interval before inheriting any architectural mutation.