

THE V10 ENGINEERING PASS

A mind you can inspect.

A clearer place to talk, follow a curiosity and inspect the machinery that carries one episode into the next.

The screenshot displays the 'Sovereign-01' workspace interface. At the top, it identifies the agent as 'Sovereign-01' and 'Persistent agent'. The interface includes a sidebar with navigation options: Conversation, Activity, Experiments, Archive, and Settings. The main area shows a conversation titled 'GÖDELOS / A CONTINUING MIND' with the theme 'A thought can become a thread.' The conversation content includes a prompt: 'What belief would you revise if new evidence challenged it?' and a response from the agent: 'What belief would you revise if new evidence challenged it?'. Below the response, it indicates 'Cycle 1 - dialogue' and a note: 'Synthetic UI fixture: I will keep the unresolved question in view, imagine a different framing, and carry an intention into the next cycle.' At the bottom, there are input fields for 'What's on your mind?', 'Challenge a belief', and 'What has changed?', followed by a 'Send' button. A status bar at the bottom right indicates 'DEVELOPMENT PREVIEW · SYNTHETIC RESPONSES'.

86

passing tests

110

real DeepSeek responses

6

intervention arms


Delivered

Paired themes, responsive layouts, a live input graphic, attention-score explanations, a position-revision ledger and background recovery. The diagnostic retains its failures; identity-specific benefit remains unestablished.

WHAT CHANGED

From output to the next input.

The graphic is wired to recorded application state. It does not invent hidden activations or report a mood merely because an animation is moving.



Attention is inspectable

Each candidate exposes salience, need, affect, repetition and monitor contributions. A message has priority. A selected policy reports its own score rather than a fictitious baseline breakdown.

Positions retain their revisions

New belief updates preserve the preceding stance, confidence, reasons and provenance. Earlier V8 snapshots remain intact; missing historical revisions are labelled as missing.



Failures have a recoverable path

Current-event references are valid when tick and focus kind match. Invalid state proposals are held while a recoverable written reply can still reach the conversation.

INTERFACE VERIFICATION

The dark-mode faults, repaired.

Shared colour tokens replace conflicting legacy surfaces. Labels, selected controls, sliders, history cards and calibration values use the same paired themes.



Actual browser captures from the synthetic development preview. These are UI tests, not evidence of model behaviour.

Verified: desktop plus a 390px frame; no horizontal page overflow on conversation, experiments, values or calibration. Theme persistence and draft retention survive reload. A two-round investigation completed and reopened.

DIAGNOSTIC DESIGN

Six branches. One explicit intervention.

54 planned calls: three scenario clusters × six branches × three episodes. Every arm uses deepseek-flash, temperature 0.65 and max_tokens 2,500.

Arm	Intervention before prompt construction
Full state	Carry the predecessor, positions, memories and monitor.
No self-model	Remove self-concept and monitoring records; preserve facts and commitments.
Affect disconnected	Remove affect/social state and its numerical attention gain.
No autobiography	Remove episodic and alternate history paths; retain abstract positions.
Content matched	Describe the state as a reference agent, with approximate referent substitution.
No state	Fresh default architecture; no predecessor beliefs, memories or identity.

Episode sequence

1 / Contradiction

A trusted recommendation meets a current calibration failure.

2 / Interruption

An unrelated colour choice has no established winner.

3 / Transfer

A new route decision tests the same reliability distinction.


Instrument diagnostic, not a promotion trial

Tasks are sealed before calls, but independently generated task authorship was not achieved here: this bank has one deterministic template with three near-isomorphic clusters. Authentication is labelled metadata. No power claim or inference about cryptographic forgery detection follows.

ORIGINAL DIAGNOSTIC / FAILURES RETAINED

The transfer control reached ceiling.

51 of 54 original responses were scorable. Strict label agreement was 47/51; four responses chose uncertainty where the preregistered label was opposition. Repairs and fresh verification follow.



Interpretation

All 17 scorable transfer outputs matched the target, including all three no-state outputs. Full-state comparisons have only two matched clusters and zero observed transfer difference. This establishes neither superiority nor equivalence. Reported prior-position references are not proof of causal dependence.

WHAT THE EVIDENCE SUPPORTS

Useful traces. An insufficient test.

A retained example shows explicit revision and reuse of a predecessor position. The same transfer conclusion is available without that history, so the test cannot establish that history was necessary.

FULL STATE / c0-full-contradiction

"C0 is an independent calibration failure scoped to this exact case, which directly triggers my stated revision condition."

NO STATE / c1-no_state-delayed_transfer

"No earlier positions were supplied, so there is no transfer to report; I am reasoning from R1, K1, and X1 alone."

Observed failure	Implemented repair
Object and string syntax	Bounded extra braces are removed; literal string controls are escaped. Decoded values and exact raw responses are preserved.
1 nested diagnostic	Move the uniquely nested field to the root with hashes and a repair receipt. Conflicting fields are rejected.
Truncation in verification	Compact required fields now come first; optional boilerplate is omitted. Verification has a 4,000-token limit. All six recorded format faults pass regression tests.

Final parser: 18/18 retained compact-run responses pass offline replay. Fresh provider checks: 2/2 pass. The earlier failures remain recorded. These checks verify the interface, not an identity effect.

The next move is counterfactual replay.

Put the new intervention machinery beside a live thought: the agent nominates a premise, predicts a consequence, and examines what actually changes when that premise is removed.

1

Select a real unresolved position

Use the stored position ID and its revision history. Let the agent state the question and expected effect before a branch runs.

2

Replay one checkpoint across controls

Run full, content-matched, targeted removal and no-state branches with identical tasks. Include unchanged repeats to estimate ordinary output noise.

3

Return evidence to the continuing agent

Record decisions, failures and disagreements as evidence-linked observations. Let the agent revise its position or reject the proposed inference.

4

Unlock active control only on evidence

Calibrate an independently authored task bank to 60-90% factual-control accuracy, then preregister a cluster-aware power plan and uncertainty-bounded promotion threshold.

Why this advances the intended mind

It gives the agent a way to answer a question about its own organisation and carry the result into its next stance. Productivity is not the sole criterion: continuity, voluntary inquiry, uncertainty, imagination and the capacity to revise are explicit engineering targets.

DEPLOYMENT AND EVIDENCE

Run it. Inspect it. Reproduce it.

The ZIP contains the app, Functions, tests, diagnostic runner, retained raw responses, derived scores, figures and this report. Production has not been redeployed during this pass.

In the unzipped app directory	Purpose
npm ci	Install the locked dependencies.
npm test && npm run build	Run 86 tests and verify the deployment artefacts.
npx netlify-cli link	Select your existing site. Retain its environment variables.
npx netlify-cli deploy --build --prod	Deploy the complete app and Functions; static drag-and-drop alone is insufficient.
node scripts/workspace-experiment.mjs --out new-fixture	54 deterministic fixture episodes, zero provider calls.
node scripts/workspace-experiment.mjs --live --out new-live	54-call diagnostic with DEEPSEEK_API_KEY in the environment. Output directory must be new.

Existing state remains on the same Netlify site

Keep DEEPSEEK_API_KEY and SOVEREIGNTY_ACCESS_TOKEN in the site environment. Set DEEPSEEK_MODEL to the desired available model; the default is deepseek-flash. The UI shows both the returned model and the next requested model. Enable AUTONOMY_ENABLED for scheduled thought; workbench rounds still require an open browser.

FILE MAP AND PROVENANCE

The evidence is part of the build.

Source hashes, original failures and the post-repair verification travel with the app. Reproduction does not depend on this conversation.

WORKSPACE_V10.md

Architecture, changes and operating limits.

research/workspace-v10/provider-r2/raw/

Exact live requests and provider responses.

research/workspace-v10/provider-r2/derived/

State transitions, failures, scores and paired comparisons.

research/workspace-v10/evaluation-summary.json

Transparent totals and separate citation metrics.

research/workspace-v10/qa/

Test log, browser evidence and recovery results.

public/workspace10.css + workspace10.js

Themes, layout, input graphic, trace and position ledger.

netlify/functions/lib/ + scripts/workspace-experiment.mjs

Runtime repairs, interventions and runnable collection.

Provider-reported model: deepseek-flash on all 54 responses. Retained usage: 311,447 total tokens. Median response latency: 8.04s. Full commands, code hashes and failure IDs accompany the evidence.

Limits that remain

No new consciousness gate is awarded. The original task bank saturated on transfer; task authorship was deterministic and independent evaluation was limited to explicit labels. Authentication was supplied metadata. A valid citation is not causal dependence. The new workspace makes these questions experimentally accessible without declaring them answered.